Name = ~~[illegible]~~ Aditya Singh

Roll No 2301010229

Date / /

Page No.

Assignment 4

1. Distributed Deadlock Detection simulation

→ Fragments :-

- S_1 : $P_1 \rightarrow P_2$ $P_3 \rightarrow P_4$
- S_2 : $P_2 \rightarrow P_5$, $P_5 \rightarrow P_6$
- S_3 : $P_6 \rightarrow P_1$

(a) Global wait for graph (combined

$P_1 \rightarrow P_2 \rightarrow P_5 \rightarrow P_6 \rightarrow P_1$ (cycle)

Also, $P_3 \rightarrow P_4$ (separate)

(b) deadlock ?

→ Yes, Processes involved in deadlock cycle ; P_1, P_2, P_5, P_6

c suggested distributed algorithms

→ use the chandy → mishra → Hass edge -chasing (probe) Algorithm for distributed deadlock detection each sits sends probes along wait for edges to detect cycles without centralized graph assembly.

2 Distributed file system performances.

Given: local = 5ms , remote = 25 ms

prob(remote) = 0.3

(a) Expected across time :-

$E = 0.7 \times 5ms + 0.3 \times 25ms$

$= 3.5 + 7.5$

$= 11 ms$

Date / /
Page No.

(b) caching strategy :-
client side read cache with LRU + TTL based validation

→ Justification :- Frequently-read remote files will be served locally reducing remote access (0.3 fraction), LRU evicts less used items. TTL keeps staleness bounded. Improves average latency while keeping consistency manageable

3 checkpointing mix to meet RPO = 1s

→ Given)
full = 200ms, incremental = 80ms, RPO = 1s

(a) proposed mix (over 10s) :-
- Take one full checkpoint every 10s (at t = 0 in period)
- Take incremental checkpoints every 1s (at t = 1, 2 - - - 9)

Total overhead (per 10s) : 1 × 200ms + 9 × 80ms = 200 + 480 => 680ms

(b) reasoning :-
- with incremental 1s, the maximum work lost on failure ≤ 1s → meets RPO.
- full once per 10s bound recovery time
- This mix minimizes full checkpoint cost while keeping incremental frequency high enough to meet RPO

Date / /
Page No.

9 case study - Global e-commerce platform:-

(a) Distributed scheduling challenges for flash sales:-

- massive, sudden split in request
- Geographic distribution & latency - data locality matters for latency and inventory correctness
- Heterogeneous nodes
- Stateful server

→ suggest algorithms for load balancing

→ Hybrid Approach

(b) Fault - tolerance strategy (RTO & RPO):-

- Active :- active multi-region deployment : service run (or) currently in multiple regions so failurous is seamless (RTO ≈ near - zero at service level)

- Data strategy

→ critical transactional data :- use synchron replication within the region to guarantee consistency and low RPO, cross region replication can be asyn chronous but with frequent replication to keep RPO small.

→ catalog / less critical data :- use multiple region eventual consistency with frequent asyn replication and cache

Date / /
Page No.

- operational measures :- chaos testing, backups runbook, and automata scaling to maximize RTO, use snapstick + incremental backup to bond RPO